

# CISC 5597 / 6935: Distributed Systems

Lecture 9: Concurrency Control

# Multiple-objects, multiple-servers

- Consistency for multiple-objects, multiple-servers
  - Single Server Case
    - Two Phase Locking
  - Distributed Transactions
    - Two Phase Commit

# Concurrency Control

- In the context of concurrency, what have we been focused on avoiding?
  - Generally so far, we have been focused on various ways to ensure safe concurrent access to a single resource.
  - Unsafe concurrent access to shared resources.
    - e.g., race proposals in Paxos.
  - Going forward, we want to handle composite operations— we want to build concurrent systems which act on multiple distinct objects

# Composite Operations

- Let's start with an example of composite operations.
- Imagine a banking application with 3 functions that can access account information:



- If each of these 3 functions were implemented in a thread-safe way, we should be OK, right?
  - NO

# Composite Operations

- Multi-threading works?

**Assuming these 2 threads are concurrently executing... What can happen?**

```
// thread 1: transfer 100  
// from savings->current  
withdraw(savings, 100)  
deposit(current, 100)
```

```
// thread 2: check balance  
s = getBalance(savings)  
c = getBalance(current)  
tot = s + c
```

# Problems with Composite Operations

- Two issues
  - Insufficient Isolation
    - Individual operations each being atomic is not enough
    - Instead, we want the deposit & withdraw making up the transfer to happen as one atomic operation...
  - Fault Tolerance
    - In the real-word, programs (or systems) can fail
    - Need to make sure we can recover safely if failure happens (particularly– no corrupted data!)

# Transactions!



*Yep! The database kind.*


# What is Transaction?

- A set of operations that need to be performed together.

- Example 1: transferring money between accounts
- Example 2: shopping cart checkout

- Initially Theo and Rodrigo have \$100.
- Goal: Transfer \$50 from Rodrigo to Theo.

Read(R)  
Update(R, \$50)  
Read(T)  
Update(T, \$150)



**Ideal: either the whole 4 operations happen, or none happen Worst case: only a subset occur**

# Transaction Background


- Transactions are defined as having the 4 ACID properties.
  - **Atomicity:** Each transaction completes in its entirety or is aborted. If aborted, should not have effect on the shared global state.
  - Example: Update account balance on multiple servers
- **Consistency:** Each transaction preserves a set of invariants about global state. (Nature of invariants is system dependent).
- Example: in a bank system, law of conservation of \$\$
- **Isolation:** Also means serializability. Each transaction executes as if it were the only one with the ability to read/write shared global state.
- **Durability:** Once a transaction has been completed, or “committed” there is no going back. In other words, there is no “undo”.

# Distributed Transactions Semantics

- Transaction Manager (TM)
  - Server in charge of orchestrating the transaction
- Steps for transaction
  - Client initiates a transaction
  - TM gives client a Transaction ID
  - Client submits operations to TM
  - TM relays operations to replicas
  - Client commits transaction
  - TM performs two phase commit

## Transaction Manager

- TM hands out TIDs
- TM manages and relays operations to Replica Leaders
- TM keeps track of all Replicas involved in the transactions



## Replica Manager

- Prepare operations
- Store operations locally in log
- But **do not commit** operations

Node C

Node F

## Client-Side Code:

```
tid = openTransaction();
RVal = a.get(tid, Rodrigo);
a.update(tid, Rodrigo, RVal - 50);
Tval = b.get(tid, Theo);
b.update(tid, Theo, Tval + 50);
closeTransaction(tid) or
abortTransaction(tid)
```

# A Transaction Example: Bank

- Array  $\text{Bal}[i]$  stores balance of Account “ $i$ ”
- **Goal:** Implement:  $\text{xfer}$ ,  $\text{withdraw}$ ,  $\text{deposit}$

```
xfer(i, j, v):
    if withdraw(i, v):
        deposit(j, v)
    else
        abort
```

```
deposit(j, v):
    Bal[j] += v
```

```
withdraw(i, v):
    b = Bal[i]          // Read
    if b >= v          // Test
        Bal[i] = b-v   // Write
        return true
    else
        return false
```

# A Transaction Example: Bank

- Imagine the following scenario:
  - $\text{Bal}[x] = 100, \text{Bal}[y] = \text{Bal}[z] = 0$
- 2 transactions:
  - T1:  $\text{xfer}(x,y,60)$
  - T2:  $\text{xfer}(x,z,70)$
- **What will happen if these 2 transactions follow ACID properties and happen in some serial order?**

# A Transaction Example: Bank

- T1; T2:
  - T1 Succeeds;
  - T2 Fails.
  - $\text{Bal}[x]=40, \text{Bal}[y]=60$
- T2; T1:
  - T2 Succeeds.
  - T1 Fails.
  - $\text{Bal}[x]=30, \text{Bal}[z]=70$
- What if we weren't careful about ACID properties? Is there a race condition?
  - ACID VIOLATION: NOT ISOLATED, NOT DURABLE
  - Updating  $\text{Bal}[x]$  with Read/Write interleaving of T1,T2
  - $\text{Bal}[x] = 30 \text{ or } 40; \text{Bal}[y] = 60; \text{Bal}[z] = 70$

# A Transaction Example: Bank

- **Goal:** Implement: xfer, withdraw, deposit
- For Consistency, we can implement **sumbalance()**
  - **Do you see anything wrong with this?**
  - **State invariant sumbalance=100 violated!**

```
xfer(i, j, v):
    if withdraw(i, v):
        deposit(j, v)
    else
        abort
```

```
deposit(j, v):
    Bal[j] += v
```

```
withdraw(i, v):
    b = Bal[i]          // Read
    if b >= v          // Test
        Bal[i] = b-v   // Write
        return true
    else
        return false
```

```
sumbalance(i, j, k):
    return Bal[i]+Bal[j]+Bal[k]
```

sumbalance()  
    ↖

# Implementing Transaction with Locks

- We can use locks to wrap xfer.

```
xfer(i, j, v):
    lock()
    if withdraw(i, v):
        deposit(j, v)
    else
        abort
    unlock()
```

## Solution 2:

- Grab short-term fine-grained locks on an item before access
- Lock(A) Read(A), Unlock(A), Lock(B) Write(B), Unlock(B) ....

```
xfer(i, j, v):
    lock(i)
    if withdraw(i, v):
        unlock(i)
        lock(j)
        deposit(j, v)
        unlock(j)
    else
        unlock(i)
        abort
```

- Solution 1:
  - Grab global lock before transaction starts – Release global lock after transaction commits
- However, is this a good approach?
  - Sequential bottleneck due to global lock!

## Problem Solved?

Nope. Consistency violation. Releasing lock i early before another transaction can cause consistency violation.

# Implementing Transaction with Locks

- Potential fix:
  - Release locks when update of all state variables complete.
  - Works now?
  - Nope. Deadlock!
    - $\text{Bal}[x]=\text{Bal}[y]=100$
    - $\text{xfer}(x,y,40)$  and  $\text{xfer}(y,x,30)$

```
xfer(i, j, v):  
    lock(i)  
    if withdraw(i, v):  
        lock(j)  
        deposit(j, v)  
        unlock(i)  
        unlock(j)  
    else  
        unlock(i)  
    abort
```

# Strict Two-Phase Locking (2PL)

# Problems for Solution 2 above

T1

```
x = Read(A)  
y = Read(B)  
Write(A, x-100)  
Write(B, y+100)
```

T2

```
x = Read(A)  
y= Read(B)  
Print(x+y)
```

Possible? (short-term, fine-grained locks on reads/writes)

R(A),R(B), W(A), R(A),R(B) C W(B) C

Locks on writes should be held  
till end of transaction

Read an uncommitted value

# Solution 3

- Solution 3
  - fine-grained locks
  - **long-term locks for writes**
    - grab lock before write, release lock after tx commits/aborts
  - **short-term locks for reads**

# Solution 3's problem

T1

```
x = Read(A)  
y = Read(B)  
Write(A, x-100)  
Write(B, y+100)
```

T2

```
x = Read(A)  
y= Read(B)  
Print(x+y)
```

Possible? long-term locks for writes, short-term locks for reads  
 $R(A), R(B), W(A), \textcolor{magenta}{R(A), R(B)}, C$   $W(B), C$

Read locks must be held  
till commit time

Non-repeatable reads

$\textcolor{magenta}{R(A)}, \textcolor{blue}{R(A), R(B)}, W(A), W(B), C$   $\textcolor{magenta}{R(B)}$

# Transactions: Split Into 2 Phases

- For reliability, transactions are split into two phases:
- **Phase 1 (Growing): Preparation:**
  - Determine what has to be done, how it will change state, without actually altering it.
  - Generate Lock set “L”
  - Generate List of Updates “U”
- **Phase 2 (Shrinking): Commit or Abort:**
  - Everything OK, then update global state
  - Transaction cannot be completed, leave global state as is
  - In either case, RELEASE ALL LOCKS

# Strict 2 phase locking (2PL)

- 2 phase locking (2PL)
  - A growing phase in which the transaction is acquiring locks
  - A shrinking phase in which locks are released
- In practice,
  - The growing phase is the entire transaction
  - The shrinking phase is at the commit time
- Optimization:
  - Use read/write locks instead of exclusive locks
  - **Read lock (Shared lock):** others can still read it
  - **Write lock (Exclusive lock):** others CANNOT access to it

# Strict Two-phase Locking (Strict 2PL) Protocol:

## TXNs obtain:

- An **X (*exclusive*) lock** on object before **writing**.
  - If a TXN holds, no other TXN can get a lock (S or X) on that object.
- An **S (*shared*) lock** on object before **reading**
  - If a TXN holds, no other TXN can get *an X lock* on that object
- All locks held by a TXN are released when TXN completes.

Note: Terminology here- “exclusive”, “shared”- meant to be intuitive- no tricks!

These policies ensure that no conflicts (RW/WR/WW) occur!


# 2PL in practice: an example

```
RLock(A)
x = Read(A)
RLock(B)
y = Read(B)
WLock(A)
buffer A←x-100
WLock(B)
buffer B←y+100
T1 issues commit :
log (A←0, B←200)
Write(A, 0)
Unlock(A)
Write(B, 200)
Unlock(B)
Commit returns
```

```
RLock(A)
x = Read(A)
RLock(B)
y = Read(B)
Print(x+y)
Unlock(A)
Unlock(B)
```

# Example

A schedule that is **not** conflict serializable:



No way for the actions of conflicts A & B to *both* happen in this order in a serial schedule!


# Conflict Dependency Graph

- Node for each committed TXN  $T_1 \dots T_N$
- Edge from  $T_i \rightarrow T_j$  if an actions in  $T_i$  **precedes** and **conflicts** with an action in  $T_j$


Theorem: Schedule is **conflict serializable** if and only if its dependency graph is acyclic

# Conflict Dependency Graph

Example:



Conflict dependency graph:




A non-conflict serializable schedule has a **cyclic** conflict dependency graph!

# More on 2PL

- What if a lock is unavailable?
  - Wait
  - Wait-for graph
  - Deadlocks possible?
- How to cope with deadlock?
  - Grab locks in order? No always possible
  - Transaction manager detects deadlock cycles and aborts affected transactions
  - Alternative: timeout and abort yourself wait detect & abort

# Deadlock Detection

Example:



In general, must search through this big graph. Sounds expensive! Is it?



Deadlock!

# Distributed Transactions

## --- 2 Phase Commit

# Distributed Transactions


- Partition databases across multiple machines for scalability
  - E.g., machine 1 responsible for account m, machine 2 responsible for account n



- Transactions often touch more than one data partition.
- How do we guarantee that all of the partitions commit the transactions, or none commit the transactions?
  - Transferring money from m to n.
  - Requirement: both machines do it, or neither.

# Modeling Distributed Transactions

- Multiple servers ( $S_1, S_2, S_3, \dots$ ), each holding some objects which can be read and written within client transactions
- Multiple concurrent clients ( $C_1, C_2, \dots$ ) who perform transactions that interact with one or more servers
  - E.g.  $T_1$  reads  $x, z$  from  $S_1$ , writes  $a$  on  $S_2$ , reads + writes  $j$  on  $S_3$
  - E.g.  $T_2$  reads  $i, j$  from  $S_3$ , then writes  $z$  on  $S_1$
- A successful commit implies agreement at all servers




# Enabling Distributed Transactions

- Similar idea as before, but:
  - State spread across servers
  - Single transaction should be able to read/modify global state while maintaining ACID properties.
  - Failures!
- Overall Idea:
  - Client initiates transaction.
  - Makes use of “coordinator”
  - All other relevant servers operate as “participants”
  - Coordinator assigns unique transaction ID (TxID)
- **Distributed Commit**
  - 2-phase commit protocol (2PC)

# Distributed Commit

- A simple distributed commit protocol:
  - Coordinator tells other processes that are involved (called participants) whether or not to locally perform the operation in question.
  - This is called the one-phase commit protocol.
- Even without failures, a lot can go wrong, for example:
  - Account j on Server 2 has only \$90
  - Account m doesn't exist!



# 2-Phase Commit

- **Phase 1: Prepare & Vote**


- Participants figure out all state changes
- Each determines if it can complete the transaction
- Communicate with coordinator

- **Phase 2: Commit**

- Coordinator broadcasts to participants: COMMIT / ABORT
- If COMMIT, participants make respective state changes


# Implementing 2-Phase Commit

- Implemented as a set of messages.




# Implementing 2-Phase Commit

- Messages in first phase
  - 1A: Coordinator sends **CanCommit?** Message to participants
  - 1B Participants respond: **VoteCommit** or **VoteAbort**



# Implementing 2-Phase Commit

- Messages in first phase
  - 1A: Coordinator sends **CanCommit?** Message to participants
  - 1B Participants respond: **VoteCommit** or **VoteAbort**
- Messages in the second phase
  - 2A: All VotedCommit: , Coordinator sends **DoCommit**
  - If any VoteAbort: abort transaction.  
Coordinator sends **DoAbort** to everyone  
=> release locks



# Timeout and Failure Cases (1)


- Coordinator times out after CanCommit?
  - Hasn't sent any commit messages, safely abort
  - Conservative. Why?
  - Preserve correctness, sacrifice performance
- Participant times out after VoteAbort
  - Can safely abort unilaterally.
  - Why?

# Timeout and Failure Cases (2)



- Participant times out after VoteCommit
  - Are unilateral decisions possible? Commit, Abort?
  - Participant could wait forever
- Solution: ask another participant (gossip protocol)
  - Learn coordinator's decision: do the same
  - Other participant hasn't voted: abort is safe. Why?
  - Coordinator has not made decision
  - No reply or other participant also VoteCommit: wait
  - 2PC is “blocking protocol” → 3PC (not covered).

# 2-Phase Commit


- Phase 1:
  - Coordinator sends request for votes
  - Participants vote
- Phase 2:
  - Coordinator counts votes
  - Coordinator informs of transaction status
- At the end of Phase 2: either all participants commit or all abort





# Coordinator State Diagram



# Participant State Diagram



# State Diagrams for Two Phase Commit




# Two Phase Commit With Failures


- What is the impact of failures on 2PC?
- 2PC is synchronous
  - Failure == node failure or network failure
  - Failure --> the protocol blocks/stalls

Coordinator fails, participants will be uncertain (waiting)


Participant fails, coordinator will be waiting



# Crash Points




**Coordinator**




**Participant**

# Two Phase Commit With Failures

- What is the impact of failures on 2PC?
- 2PC is synchronous
  - Failure == node failure or network failure
  - Failure --> the protocol blocks/stalls
- Detect failure using Timeouts
  - Coordinator detects participant failure and assumes ABORT à Transaction terminates
  - Participant detects Coordinator failure
  - Why Can't Participant automatically ABORT?




# Participant Recovery from Coordinator Failure



- Participant in **Uncertain** state
  - waiting for coordinator to say “commit” or “abort”
- It detects failure of coordinator
  - Using timeout
- Participant in **Uncertain** state
  - can't assume either outcome
    - BAD things happen if participant makes wrong assumptions
  - waits for coordinator to restart
    - On restart contact coordinator for final outcome


# Coordinator Recovery from Participant Failure

- Coordinator in wait state
  - waiting for participant to say “commit” or “abort”
- It detects failure of participant
  - Using timeout
- Coordinator assumes they said ‘no’
  - Takes no response as an abort
  - Abort transaction!!
- If participant Fails, Coordinator can make progress



# Two Phase Commit With Failures

- Detect failure using Timeouts
  - Coordinator detects participant failure and assumes ABORT a transaction terminates
  - Participant detects Coordinator failure
  - The participant must wait for coordinator
  - The transaction is stalled!



# Properties of 2-Phase Commit

- Correctness
  - Neither can commit unless both agreed to commit
- Performance
  - $3N$  messages per transaction
- How to handle failure?
  - Timeouts à performance bad in case of failure!

# Deadlocks and Livelocks

- Distributed deadlock
- Cyclic dependency of locks by transactions across servers
  - Q: How can this happen?
- In 2PC this can happen if participants unable to respond to voting request
  - e.g. still waiting on a lock on its local resource.
  - Q: How can we fix this?
- Handled with a timeout. Participants times out, then votes to abort. Retry transaction again.
  - Addresses the deadlock concern
  - However, danger of LIVELOCK – keep trying!

# Summary

- Distributed consistency management
- ACID Properties desirable
- Single Server case: use locks + 2-phase locking (**strict 2PL**), transactional support for locks
- Multiple server distributed case: use **2-phase commit** for distributed transactions. Need a coordinator to manage messages from participants
- 2PC can become a performance bottleneck